

Madhumitha Sekamuri

madhutrish06@gmail.com | +1-614-483-7876 | [LinkedIn](#) | [GitHub](#) | Ashburn, VA

Summary

Data Scientist with proven experience delivering scalable ML, anomaly detection, and LLM driven solutions across healthcare, insurance, and enterprise data systems. Adept in predictive modeling, generative AI, and cloud platforms (AWS, Azure, GCP). Skilled in the full ML lifecycle from EDA and feature engineering to deployment and optimization. Passionate about combining GenAI, anomaly detection, and knowledge graphs to deliver explainable insights in high-stakes domains including audit, healthcare, and risk analysis.

Experience

Moxie

Ashburn, VA (remote)

Data Scientist

Feb 2025 – Present

- Developed and deployed a Retrieval-Augmented Generation (RAG) pipeline using LangChain, FAISS, and OpenAI APIs on AWS Lambda and API Gateway, enabling source-grounded summarization of long-form audit and clinical documents; improved document review efficiency by 60% while ensuring explainability for regulatory compliance.
- Engineered a domain-specific LLM-powered chatbot leveraging ChromaDB, OpenAI embeddings, and FastAPI, hosted on Amazon ECS, to support HR/IT knowledge base queries; integrated prompt engineering and context-aware retrieval, achieving 90% response relevance and reducing manual inquiry resolution time.
- Implemented an anomaly detection and entity extraction framework with GPT-4, Pydantic, and Docker containers deployed on AWS Fargate, automating the identification of irregular records, extracting structured entities (medications, symptoms, diagnoses), and providing audit ready traceability for downstream analytics.

Behavioral Neuroscience – The Ohio State University

Columbus, OH

Data Scientist

May 2024 – Dec 2024

- Implemented a SLEAP-based multi-animal tracking system using Python, TensorFlow, and OpenCV on Google Cloud Compute Engine, enabling scalable processing of high-frame-rate rodent behavior videos and improving annotation efficiency by 70%.
- Built a deep learning pipeline integrating CNNs and LSTM architectures in PyTorch, trained on Google Cloud AI Platform, for pose estimation and sequential behavior classification, achieving 92% accuracy on labeled mouse interaction data.
- Applied transfer learning with a U-Net model using Keras and TensorFlow, trained on GCP TPUs, to detect fine-grained social behaviors (Ano-genital sniffing, nose-to-nose contact), achieving an F1-score of 0.92 on custom datasets.
- Built baseline vs. fine-tuned model comparisons for pose estimation pipelines, helping quantify the added value of deep architectures (CNNs, LSTMs) over traditional models.
- Designed an automated behavioral analysis workflow using scikit-learn, ffmpeg, and NumPy, executed via Cloud Composer to streamline preprocessing, inference, and post-processing in a reproducible and scalable environment.

Pelotonia Research Center – Cancer Research

Columbus, OH

Data Science Intern

Jan 2023 – May 2024

- Built and automated end-to-end ETL pipelines using Python, SQL, Apache Airflow, and Pandas on Azure Data Factory, to ingest, clean, and integrate large-scale structured and unstructured EHR and genomic datasets, improving data accessibility and reporting efficiency by 50%.
- Developed medical image registration pipelines using MONAI, PyTorch, and SimpleITK, deployed on Azure Machine Learning Compute, applying diffusion models, affine, and deformable transformations to achieve 97% alignment accuracy across histopathology slides.
- Developed classification and survival analysis models using scikit-learn, XGBoost, and Lifelines on Azure ML Studio to predict patient outcomes and treatment response; performed Kaplan-Meier analysis and log-rank tests to identify survival patterns.
- Created interactive Tableau dashboards hosted via Azure App Service, to visualize survival curves, treatment efficacy, and patient stratification, enabling researchers and clinicians to explore model outputs and key clinical insights.

Cognizant Technology Services

Chennai, India

Machine Learning Engineer

Mar 2021- Dec 2022

- Built and deployed risk prediction models using LightGBM and pandas on Amazon SageMaker to classify high-risk insurance policies, enhancing fraud detection and underwriting decisions with a 30% improvement in accuracy across 100K+ policy records.
- Deployed ML models as REST APIs using FastAPI, containerized with Docker, and hosted on AWS ECS, while managing experiment tracking and version control via MLflow on S3, streamlining the MLOps lifecycle in a regulated insurance environment.
- Designed and orchestrated scalable data pipelines using Databricks (PySpark) and dbt, integrated with AWS Glue and Redshift, automating ingestion and transformation of 1M+ insurance records and reducing model retraining time by 40%, supporting real-time policy risk scoring workflows.
- Conducted comparative benchmarking of machine learning algorithms (LightGBM, XGBoost, Logistic Regression) using cross-validation and AUC/F1 metrics to determine optimal model for policy risk classification.
- Built and automated end-to-end ETL pipelines using Python, SQL, Apache Airflow, and Pandas on Azure Data Factory, to ingest, clean, and integrate large-scale structured and unstructured EHR and genomic datasets, improving data accessibility and reporting efficiency by 50%.

Skills

- Languages:** Python, R, Java, JavaScript, C, D3.js, HTML/CSS
- ML & DL:** Scikit-learn, XGBoost, LightGBM, Lifelines, H2O, PyTorch, TensorFlow, Keras, MONAI, CNNs, LSTMs, U-Net, LBPH, Transfer Learning, PyTorch Geometric
- LLM & NLP:** OpenAI (GPT-4, ChatGPT APIs), Google PaLM, LLaMA 2, Mistral, LangChain, Retrieval-Augmented Generation (RAG), Prompt Engineering, ChromaDB, FAISS, Pydantic, TextBlob
- Data Analytics & Visualization:** Pandas, NumPy, Matplotlib, Seaborn, Plotly, Tableau, Metabase, Exploratory Data Analysis (EDA), Kaplan-Meier Analysis, Log-rank Test
- Data Engineering & Pipelines:** Apache Airflow, dbt, SQL, MySQL, Azure Data Factory, AWS Glue, ETL Pipelines, Data Cleaning, Data Integration, Data Modeling
- MLOps & Model Deployment:** FastAPI, Flask, Docker, MLflow, REST APIs, AWS Lambda, AWS ECS, AWS Fargate, Amazon SageMaker, Azure ML Studio, Azure ML Compute, Google Cloud AI Platform, Cloud Composer
- Computer Vision & Image Processing:** OpenCV, SimpleITK, Dlib, Haar Cascades, FaceNet, Image Registration (Affine, B-spline, Deformable, Diffusion Models), High-Frame-Rate Video Analysis, SLEAP
- Cloud Platforms:** AWS (Lambda, ECS, Fargate, S3, SageMaker, Glue, Redshift), Azure (App Service, ML Studio, Data Factory, ML Compute), GCP (AI Platform, Composer, Compute Engine, TPU)
- Databases & Storage:** MySQL, PostgreSQL, Oracle SQL, SQL Server, NoSQL, DynamoDB, Redshift, Amazon S3, Elasticsearch, Neo4j
- Testing & Optimization:** Cross-Validation, Hyperparameter Tuning (Grid Search, Random Search), Model Benchmarking, F1-Score, Accuracy, Precision
- Frameworks/Libraries:** Flask, FastAPI, Node.js, Express.js, React.js, Angular.js, Apache Spark, MLlib, Airflow, Hadoop
- Data Governance & Privacy:** HIPAA compliance, PII handling, regulated data workflows (Insurance & Healthcare)
- Certifications:** Microsoft Azure Cloud Fundamentals, Azure AI Fundamentals, Google – Python for Data Science Certification, CyberArk Trustee Certification, NPTEL – Java, Machine Learning

Projects

AI-Powered Business Idea Evaluation System:

Built an end-to-end pipeline to evaluate 500+ business ideas using LLMs (LLaMA 2, Mistral, ChatGPT APIs) and classical ML models; performed text preprocessing, EDA, and model benchmarking with Logistic Regression and Decision Trees, achieving over 85% agreement with expert-labeled criteria through cross-validation and hyperparameter tuning.

Text-to-SQL App:

Built a Streamlit application that translates natural language into SQL queries using LangChain and Google PaLM, enabling seamless interaction with MySQL databases. Integrated prompt chaining, query execution, and result visualization to support real-time data exploration for business users.

Smart Mobility - Smart-Mobility-GNN

Built a GNN-based model using PyTorch Geometric on the Open Traffic Dataset, modeling road networks as graphs for route optimization and traffic prediction; achieved a 15% improvement in path prediction accuracy over baseline models.

Face-X - Face Recognition

Developed a high-accuracy facial recognition attendance system using OpenCV, Haar Cascade, and LBPH, achieving 98% accuracy in real-time identification; enhanced with Dlib and FaceNet for robust feature extraction. Presented and Published at ICIVC 2022.

Education

Master of Science in Computer Science and Engineering – Ohio State University – CGPA: 3.53

Dec 2024

Coursework: Neural Networks, Data Mining, Fairness in Artificial Intelligence and Databases, Data Visualization, Parallel Computing, Advanced OS, Cybersecurity

Bachelor of Technology in Computer Science and Engineering – Anna University – CGPA: 4

June 2021

Coursework: Problem-Solving, Object-Oriented Design, Data Structures & Algorithms, Databases, Operating Systems, System Design, Networking, Machine Learning